

NIST 800-53 as Code: Quantifying the Compliance-Exposure Reduction of Continuous Automated Evidence Collection and Its Automatability Ceiling

Harshavardhan Malla Independent Researcher
Email: harshavardhanmalla75@gmail.com

Abstract: Expressing NIST 800-53 controls as executable policy-as-code lets a federal system collect control evidence continuously, but the field has reported only that drift is caught faster, never how much compliance exposure the executable form actually removes or what caps it. We give the first quantified answer: detection latency and exposure reduction are distinct payoffs, the former collapses by two orders of magnitude while the latter is hard-capped by the share of controls that can be expressed as code, and we measure that cap. The operational payoff is a sizing rule, exposure reduction tracks the automatable-exposure share rather than detection latency, so a program can be budgeted against the catalog instead of against collection speed. Federal systems must show that their security controls remain effective over time, not only at an annual audit. Continuous monitoring guidance and machine-readable control formats make it possible to collect control evidence automatically and continuously, so a control that drifts out of compliance is caught in days rather than at the next assessment. The operational question is how large that benefit is, and what bounds it. We answer with a pre-registered simulation of a 200-control catalog drifting as a Poisson process over two years, with hypotheses, thresholds, and evaluation seeds locked before any result was inspected. On automatable controls, continuous monitoring cuts the mean time to detect drift from 272 days under annual assessment to 1 day, a factor of 272 (95% bias-corrected and accelerated bootstrap interval [269, 275]). The overall exposure reduction is bounded by automatability: continuous monitoring removes 0.567 of annual exposure ([0.553, 0.581]) against a ceiling of 0.643 ([0.628, 0.660]) set by the automatable-exposure share, because the controls a machine cannot check form a floor that no detection speed can lower. The benefit shows diminishing returns as the manual cadence tightens (33,000 control-days saved over annual but 16,000 over quarterly) and concentrates in high-drift controls: the top-drift quartile alone captures 79% of the achievable reduction. We also derive these findings in closed form. A short exposure decomposition shows that the fractional reduction is bounded above by the automatable-exposure share, $R \leq C$, that the absolute benefit must diminish as the manual cadence tightens, and that a heavy-tailed drift distribution forces the benefit to concentrate, with the closed form reproducing the frozen ceiling $C = 0.643$ and realized reduction $R = 0.567$. The deployment rule that follows is concrete. Instrument the high-drift automatable controls first, size programmatic expectations to the automatable-exposure share rather than to detection latency, and treat the manual remainder as an irreducible residual.

Index Terms: continuous monitoring, NIST 800-53, compliance automation, OSCAL, policy as code, control drift, compliance exposure, pre-registration, reproducibility

A FEDERAL authorization to operate rests on a body of security controls [1], and that authorization is a statement about a moment in time. The controls are assessed, the system is judged acceptable, and an authorizing official signs. Between that signature and the next assessment, the system keeps running and its configuration keeps changing. A control drifts out of compliance when an audit log is silently disabled, a TLS certificate expires, a firewall rule is loosened for a troubleshooting session and never restored, or a baseline configuration is edited by hand. Each such drift opens a window of exposure that lasts until someone notices. Under the traditional model of periodic assessment, noticing happens at the next audit, which for many control families is a year away.

Information security continuous monitoring [2] was introduced precisely to shorten that window. Its assessment companion [3] frames continuous monitoring as an ongoing program rather than a periodic event, and the Open Security Controls Assessment Language [4] gives the controls a machine-readable representation so that evidence can be collected, checked, and compared by software. Together these enable a compliance-as-code posture: a control is expressed as a checkable rule, the rule is evaluated on a short cadence, and a drift is caught within that cadence rather than at the next periodic assessment. The policy-as-code tooling that the wider industry has converged on, from general policy engines [5] to desired-state configuration systems [6], makes the mechanics routine for the controls that can be expressed this way. A companion study takes policy-as-code further still, using it for preventive enforcement that blocks recurring control violations before they occur rather than detecting them after [21]; here we hold the detection framing fixed and measure the exposure it removes.

The promise is attractive enough that it is easy to overstate. Two facts bound it. First, not every control is machine-checkable. A control that requires a documented and approved policy, a signed agreement, or a physical-security inspection still demands human assessment; continuous collection cannot reach it, so it stays on the periodic cadence no matter how fast the automatable controls are checked. The un-automatable remainder therefore sets a floor on the achievable reduction in exposure. Second, exposure is not spread evenly across the catalog. A control that drifts twice a decade contributes little

exposure whatever the monitoring regime, while a control that drifts monthly dominates the total. The value of automation is concentrated where drift is frequent, and those same controls, the technical ones, are also the most machine-checkable.

This paper quantifies both the benefit and its bounds. Prior policy-as-code work established that a control expressed as an executable rule can be evaluated on a short cadence and that drift is then caught faster; what it never measured is how much compliance exposure that executable form actually removes, or the hard limit on that removal. Our central quantified result supplies both: continuous evaluation of policy-as-code controls cuts detection latency by a factor of 272 yet removes only 0.567 of annual exposure against a measured ceiling of 0.643, because the share of controls that can be expressed as code caps the exposure that any collection speed can remove. This separation of a latency payoff from an exposure payoff, and the measurement of the ceiling on the latter, is what prior compliance-as-code accounts did not show, and it converts the case for automation from a latency story into a budgetable sizing rule. We do not claim to measure any particular agency's posture; we build a transparent model of control drift, lock a set of hypotheses and thresholds before looking at the evaluation data, and report what the model says. The contribution is fourfold.

- We separate the *detection-latency* benefit of continuous monitoring, which is enormous on the controls it reaches, from the *exposure-reduction* benefit, which is bounded by automatability. Conflating the two is the central error in optimistic compliance-as-code business cases.
- We formalize an *automatability ceiling*: the fractional reduction in compliance exposure cannot exceed the share of baseline exposure attributable to automatable controls, and we confirm that the realized reduction sits below that ceiling.
- We show the absolute benefit exhibits *diminishing returns* against the manual baseline: the tighter the existing assessment cadence, the fewer control-days continuous monitoring adds.
- We show the benefit is *concentrated*: automating only the highest-drift quartile of controls captures the large majority of the achievable reduction, which gives a direct deployment ordering.

The methodology is deliberately conservative. Every quantitative claim is a pre-registered hypothesis evaluated on seeds that were not inspected during model development, and every interval is a bias-corrected and accelerated (BCa) bootstrap interval [7], [8] rather than a null-hypothesis significance test, following the pre-registration discipline now standard in empirical software engineering [9].

II. BACKGROUND AND RELATED WORK

A. From periodic assessment to continuous monitoring

The control catalog of NIST SP 800-53 [1] is the backbone of the federal Risk Management Framework, and FIPS 199 [10] together with its mapping guidance [11] sets the categorization that determines which controls apply. Continuous monitoring [2] reframes assessment from a point-in-time

judgment into an ongoing process, and its program-assessment companion [3] provides criteria for judging whether such a process is effective. Our work is complementary: rather than assessing a monitoring program against a checklist, we model the exposure consequences of moving a control from a periodic to a continuous cadence, and we quantify the gain.

B. Compliance as code and configuration drift

Expressing controls and configurations as code is now common practice. OSCAL [4] gives controls, assessment plans, and results a structured machine-readable form; general policy engines [5] evaluate declarative rules over configuration data; and desired-state configuration systems [6] detect and correct drift from a declared baseline. The software-engineering literature has examined the reliability of this code itself: Rahman and colleagues catalog recurring security weaknesses in infrastructure-as-code scripts [12] and characterize their prevalence across ecosystems [13]. That line of work asks whether the automation is itself correct. We ask a different and complementary question: assuming the automation works, how much compliance exposure does continuous evaluation remove, and what limits it. The autonomic-computing vision of self-managing systems [14] anticipated this shift toward continuous self-assessment; our results quantify where its value is concentrated and where it saturates.

C. Why exposure, not just latency

Vulnerability-management practice has learned to distinguish the speed of a signal from its operational value. Federal patch-management guidance [15] and the CISA Binding Operational Directives [16], [17] bind agencies to remediation timelines and to asset-visibility requirements, and breach data continues to attribute incidents to known, unremediated weaknesses [18]. The lesson, also visible in hygiene-augmented prioritization research [19], [20], is that a fast signal is only worth as much as the exposure it actually removes. We carry that lesson into compliance: detection latency on automatable controls collapses by two orders of magnitude, but the exposure reduction that matters for an authorization decision is governed by how much of the catalog the automation can reach.

III. SYSTEM MODEL

We model a single information system over a fixed horizon and ask how much compliance exposure each monitoring regime leaves behind. All quantities are synthetic with documented distributions; no operational, employer, or audit data is used.

A. Control catalog

Each seed instantiates a catalog of $N = 200$ control instances. Each control i carries three attributes.

Drift hazard λ_i , the expected number of compliance-relevant drift events per year, is drawn from a two-component mixture that reflects how real catalogs behave. A high-drift technical stratum (30% of controls: logging, certificate, configuration,

and access controls) has mean rate 6 per year; a low-drift stratum (70%: policy, documentation, and physical controls) has mean rate 0.5 per year. Technical controls change often because the systems they govern change often; policy controls change rarely.

Automatability $a_i \in \{0, 1\}$ marks whether the control can be checked by software. It is true with probability 0.85 for the high-drift technical stratum and 0.35 for the low-drift stratum, encoding the empirical correlation that the controls which drift most are also the ones most amenable to machine evaluation. The realized automatable fraction is therefore a model output, not a fixed input; across seeds it is 0.492 ([0.479, 0.504]).

Remediation time is fixed at $\tau_r = 7$ days: once a drift is detected, the control returns to compliance after a week.

B. Drift and exposure

Over a horizon of $H = 730$ days, drift events for control i arrive as a Poisson process at rate $\lambda_i/365$ per day. A drift event makes the control non-compliant from its onset until the control is detected and then remediated τ_r days later. The *compliance exposure* of a control is the total time it spends non-compliant over the horizon, measured in control-days; the exposure of the system is the sum over controls. For a drift detected with latency d , the exposure contributed is $d + \tau_r$, so total exposure decomposes as

$$E = \sum_{i=1}^N \sum_k (d_{i,k} + \tau_r), \quad (1)$$

where $d_{i,k}$ is the detection latency of the k -th drift of control i . The monitoring regime enters only through the latencies $d_{i,k}$.

C. Monitoring regimes

Under *periodic- T* assessment, every control is assessed every T days and a drift is detected at the next assessment boundary, so its latency is the time from drift onset to that boundary. We evaluate the two cadences agencies actually use: annual ($T = 365$) and quarterly ($T = 90$). Under *continuous* monitoring, an automatable control is evaluated on a one-day collection cadence, so its detection latency is at most a day, while a non-automatable control stays on the periodic cadence T . Continuous monitoring thus collapses the latency of the automatable controls toward the collection interval and leaves the rest untouched, which is exactly the structure that produces a ceiling.

D. Metrics

We report four quantities. *Mean time to detect* (MTTD) is the mean detection latency over drift events, reported separately for the automatable subset where continuous monitoring acts. *Compliance exposure* is total control-days non-compliant. The *exposure reduction* of continuous monitoring against a periodic baseline is

$$R = \frac{E_{\text{periodic}} - E_{\text{continuous}}}{E_{\text{periodic}}}. \quad (2)$$

The *automatable-exposure share* is the fraction of baseline periodic exposure attributable to automatable controls,

$$C = \frac{\sum_{i:a_i=1} E_i^{\text{periodic}}}{\sum_i E_i^{\text{periodic}}}, \quad (3)$$

which is the ceiling on R : even instantaneous detection of every automatable control cannot remove the exposure of the controls that stay on the manual cadence, so $R \leq C$ up to the residual cadence and remediation latency of the automatable controls themselves.

IV. EXPERIMENTAL DESIGN

The study is pre-registered. Hypotheses, thresholds, model constants, and the evaluation seed range were fixed in a dated protocol before any result on the evaluation seeds was inspected, so that no analytic choice could be steered by the outcome.

A. Hypotheses

H1 (detection latency). On automatable controls, continuous monitoring reduces MTTD relative to annual assessment by at least a factor of 10, with the BCa interval excluding a factor of 10.

H2 (automatability ceiling). The realized exposure reduction R does not exceed the automatable-exposure share C by more than 0.03; the un-automatable controls form a floor.

H3 (diminishing returns). The absolute control-days that continuous monitoring eliminates is smaller against a quarterly baseline than against an annual one.

H4 (drift concentration). Automating only the highest-drift quartile of controls captures at least 70% of the exposure reduction achievable by automating all automatable controls.

B. Protocol and statistics

Each hypothesis is evaluated over 25 evaluation seeds (frozen range 800 to 824). For every quantity we report the mean and a 95% BCa bootstrap interval with 10,000 resamples [7], [8]; interval non-overlap, not a p -value, is the evidentiary standard. The pre-registered failure criteria are explicit: H1 is declared null if the MTTD reduction on automatable controls falls below a factor of 5; H2 is rejected if R exceeds C by more than 0.03. No drift or automatability distribution is re-tuned to reach any threshold. The development of the model used separate seeds; the evaluation seeds were touched only once, to produce the numbers below.

V. RESULTS

Table I collects the pre-registered quantities with their intervals; all four hypotheses are supported. We take them in turn.

Detection on automatable controls is two orders of magnitude faster (H1). MTTD on the automatable subset falls from 272.0 days under annual assessment ([269.1, 275.0]) to 1.0 day under continuous monitoring, a reduction factor of 272 ([269, 275]), far above both the pre-registered factor-of-10 bar and the factor-of-5 failure threshold (Fig. 2, left). On

TABLE I
PRE-REGISTERED QUANTITIES (MEAN OVER 25 SEEDS, 95% BCa INTERVAL).

Quantity	Mean	95% BCa interval
MTTD, automatable, annual (days)	272.0	[269.1, 275.0]
MTTD, automatable, continuous (days)	1.0	[1.0, 1.0]
MTTD reduction factor	272.0	[269.1, 275.0]
Exposure reduction R vs annual	0.567	[0.553, 0.581]
Automatable-exposure share C	0.643	[0.628, 0.660]
Ceiling gap $R - C$	-0.077	[-0.079, -0.075]
Control-days saved vs annual	33,266	[32,204, 34,387]
Control-days saved vs quarterly	15,584	[15,074, 16,211]
Top-drift-quartile capture	0.786	[0.777, 0.796]

the controls it can reach, continuous monitoring converts a detection latency measured in months into one measured in a day.

The overall reduction is capped by automatability (H2). Continuous monitoring removes 0.567 of annual exposure ([0.553, 0.581]) against a ceiling of $C = 0.643$ ([0.628, 0.660]). The realized reduction sits below the ceiling, with a gap of -0.077 ([-0.079, -0.075]), so the ceiling binds and is not exceeded (Fig. 1). The reduction falls short of the ceiling by a small margin because the automatable controls still carry a one-day cadence plus a seven-day remediation, but the dominant fact is structural: the un-automatable controls form a floor of roughly 36% of baseline exposure that no detection speed can lower. A program that promises to eliminate most compliance exposure through automation is making a promise the catalog cannot keep.

The absolute benefit shows diminishing returns (H3). Continuous monitoring saves 33,266 control-days over an annual baseline ([32,204, 34,387]) but only 15,584 over a quarterly baseline ([15,074, 16,211]), a difference of about 17,700 control-days (Fig. 2, right). The fractional reduction stays near the automatable share regardless, but the absolute control-days that automation buys shrink as the manual baseline tightens, because a quarterly cadence has already removed much of the latency that continuous monitoring would otherwise eliminate. An agency already assessing quarterly should expect roughly half the absolute benefit that an agency assessing annually would see.

The benefit is concentrated in high-drift controls (H4). Automating only the highest-drift quartile of controls captures 0.786 ([0.777, 0.796]) of the reduction achievable by automating every automatable control. Because exposure is the product of drift frequency and detection latency, the controls that drift most often dominate the total, and they are also the most machine-checkable. The practical consequence is an ordering: the first quarter of the automation effort delivers nearly four-fifths of the available exposure reduction.

VI. THEORETICAL ANALYSIS

The four empirical findings are not coincidences of the chosen constants; each is the visible face of a structural relationship that the exposure decomposition makes explicit. This section derives the automatability ceiling, the diminishing-returns relation, and the drift-concentration result in closed form from the System Model, and then checks the closed




Fig. 1. The exposure reduction realized by continuous monitoring (left bar) sits below the ceiling set by the automatable-exposure share (right bar); the gap between the ceiling and unity is the floor formed by controls a machine cannot check.

Alt text: A two-bar chart on a y-axis labeled fraction of annual exposure removed, from 0 to 1. The left bar, realized continuous reduction, reaches about 0.57; the right bar, automatable-share ceiling, reaches about 0.64, both with small error bars. A shaded band above 0.64 is labeled non-automatable floor.




Fig. 2. Left: mean time to detect on automatable controls falls 272-fold from annual assessment to continuous monitoring. Right: the absolute control-days saved diminish as the manual cadence tightens from annual to quarterly.

Alt text: Two side-by-side bar charts. Left, on a logarithmic y-axis of mean time to detect in days, an annual-assessment bar at 272 days towers over a continuous-as-code bar at 1 day. Right, on a linear y-axis of control-days saved per fleet over two years, the over-annual bar reaches about 33,000 and the over-quarterly bar about 15,600, both with error bars.

form against the frozen numbers. The framing mirrors the information-and-coverage view of a fundamental limit [22]: detection speed is a channel, automatability is the coverage of that channel, and no amount of channel quality can act on the part of the catalog the channel does not reach.

A. The exposure decomposition

Write the exposure of a single drift event as the sum of the time it stays undetected and the fixed remediation it then incurs. Summing over the drifts of control i and then over controls, and replacing the per-event latency by its mean, the system exposure under a monitoring regime is

$$E = \sum_{i=1}^N \lambda_i (\bar{d}_i + \tau_r), \quad (4)$$

where λ_i is the expected number of detected drifts of control i over the horizon, $\bar{d}_i$ is its mean detection latency, and τ_r is the fixed remediation time. Equation (4) is the mean-field form of the per-event sum in Eq. (1): exposure is the product of how often a control drifts and how long each drift persists, summed across the catalog. The monitoring regime enters only through the latencies $\bar{d}_i$; the drift rates λ_i and the remediation τ_r are properties of the system and the response process, not of the detector.

B. The automatability ceiling

Partition the catalog into the automatable set $\mathcal{A} = \{i : a_i = 1\}$ and its complement. Under continuous monitoring, an automatable control is evaluated on the one-day collection cadence, so its latency falls toward the collection interval $c = 1$ day, while a non-automatable control keeps the periodic latency $\bar{d}_i^T$ it had under the baseline cadence T . Decompose the baseline (periodic) exposure into the part carried by automatable controls and the part carried by the rest,

$$E_{\text{per}} = \underbrace{\sum_{i \in \mathcal{A}} \lambda_i (\bar{d}_i^T + \tau_r)}_{E_{\mathcal{A}}^{\text{per}}} + \underbrace{\sum_{i \notin \mathcal{A}} \lambda_i (\bar{d}_i^T + \tau_r)}_{E_{\mathcal{A}}^{\text{per}}}. \quad (5)$$

Continuous monitoring changes only the first sum and leaves the second untouched, so the exposure it removes is

$$E_{\text{per}} - E_{\text{cont}} = \sum_{i \in \mathcal{A}} \lambda_i (\bar{d}_i^T - c), \quad (6)$$

because the τ_r terms cancel and the latency of each automatable control drops from $\bar{d}_i^T$ to c . The fractional reduction is therefore

$$R = \frac{E_{\text{per}} - E_{\text{cont}}}{E_{\text{per}}} = \frac{\sum_{i \in \mathcal{A}} \lambda_i (\bar{d}_i^T - c)}{E_{\text{per}}}. \quad (7)$$

Now bound the numerator. For every automatable control $\lambda_i (\bar{d}_i^T - c) \leq \lambda_i (\bar{d}_i^T + \tau_r)$, because $c \geq 0$ and $\tau_r \geq 0$, so the numerator of Eq. (7) is at most $E_{\mathcal{A}}^{\text{per}}$. Dividing by E_{per} ,

$$R \leq \frac{E_{\mathcal{A}}^{\text{per}}}{E_{\text{per}}} \equiv C, \quad (8)$$

which is exactly the automatable-exposure share of Eq. (3). Equation (8) is the *automatability ceiling*: no matter how fast or how cheap the collection cadence becomes, continuous monitoring cannot remove more than the share of baseline exposure that the automatable controls carry, because the exposure of the un-automatable remainder $E_{\mathcal{A}}^{\text{per}}$ is invariant under the detector. The floor on residual exposure is $1 - C$ of the baseline, and it is a property of the catalog, not of the monitoring program.

The bound is tight only in the unreachable limit $c \rightarrow 0$ and $\tau_r \rightarrow 0$; at the model's operating point the realized reduction sits strictly below C by the residual that the automatable controls themselves still carry. Substituting $\bar{d}_i^T - c = (\bar{d}_i^T + \tau_r) - (c + \tau_r)$ into Eq. (7) and writing ρ for the exposure-weighted mean of $(c + \tau_r)/(\bar{d}_i^T + \tau_r)$ over $\mathcal{A}$ gives the exact relation

$$R = C(1 - \rho), \quad (9)$$

so the ceiling gap $C - R = C\rho$ is the fraction of automatable exposure that the residual one-day cadence and seven-day remediation leave in place. The gap is small but strictly positive precisely because c and τ_r are small relative to the periodic latency on automatable controls, not zero.

C. Diminishing returns against the manual cadence

The absolute benefit is the numerator of Eq. (7), the control-days saved,

$$\Delta(T) = E_{\text{per}}(T) - E_{\text{cont}}(T) = \sum_{i \in \mathcal{A}} \lambda_i (\bar{d}_i^T - c). \quad (10)$$

For a Poisson drift detected at the next assessment boundary of a periodic- T regime, the mean latency is increasing in T : a longer interval between assessments leaves a drift undetected longer, so $\bar{d}_i^T$ grows with T and $\bar{d}_i^{T'} \leq \bar{d}_i^T$ whenever $T' \leq T$. Continuous monitoring drives the same automatable latencies to c regardless of T . Therefore, term by term in Eq. (10),

$$\Delta(T') = \sum_{i \in \mathcal{A}} \lambda_i (\bar{d}_i^{T'} - c) \leq \sum_{i \in \mathcal{A}} \lambda_i (\bar{d}_i^T - c) = \Delta(T) \quad (11)$$

for $T' \leq T$. The tighter the manual cadence, the smaller the absolute control-days continuous monitoring can add, because the manual cadence has already removed part of the latency that continuous monitoring would otherwise eliminate. The fractional reduction, by contrast, stays near C across cadences, because the denominator $E_{\text{per}}(T)$ shrinks with T at nearly the same rate as the numerator. Diminishing returns are therefore an absolute, not a fractional, phenomenon.

D. Drift concentration

Equation (10) is a sum weighted by λ_i , so the controls that drift most often dominate the saved exposure. Let the automatable controls be sorted by descending λ_i and let Q be the highest-drift quartile. Because the per-control saving $\lambda_i (\bar{d}_i^T - c)$ is itself proportional to λ_i (the latency factor $\bar{d}_i^T - c$ varies little across automatable controls under a common cadence), the share of $\Delta(T)$ captured by Q is approximately

$$\frac{\sum_{i \in Q} \lambda_i (\bar{d}_i^T - c)}{\sum_{i \in \mathcal{A}} \lambda_i (\bar{d}_i^T - c)} \approx \frac{\sum_{i \in Q} \lambda_i}{\sum_{i \in \mathcal{A}} \lambda_i}. \quad (12)$$

When the drift distribution is heavy-tailed, the right-hand side of Eq. (12) is large even though Q is a quarter of the controls. In the model the rate is a two-component mixture in which a 30% technical stratum drifts at 6 per year and the remaining 70% at 0.5 per year, a twelve-fold ratio. The high-drift stratum thus carries a disproportionate share of $\sum_i \lambda_i$, and because that stratum is also the most automatable, the highest-drift quartile captures the large majority of the achievable reduction. Concentration follows from heavy-tailedness alone and would survive any drift distribution with a comparable spread.

E. Numerical verification against the frozen results

The closed form must reproduce the frozen numbers, not merely point in their direction. The ceiling of Eq. (8) is the automatable-exposure share, which the frozen results report as $C = 0.6435$ ($[0.628, 0.660]$), and the realized fractional

TABLE II

SENSITIVITY TO THE MANUAL BASELINE CADENCE (MEAN OVER 25 SEEDS, 95% BCA INTERVAL). THE FRACTIONAL REDUCTION TRACKS THE AUTOMATABLE SHARE AND IS NEARLY FLAT; THE ABSOLUTE CONTROL-DAYS SAVED DIMINISH AS THE CADENCE TIGHTENS.

Manual baseline	Fractional reduction R	Control-days saved
Annual ($T = 365$)	0.567 [0.553, 0.581]	33,266 [32,204, 34,387]
Quarterly ($T = 90$)	0.552 [0.540, 0.563]	15,584 [15,074, 16,211]

reduction of Eq. (7) is $R = 0.5669$ ([0.553, 0.581]). The inequality $R \leq C$ holds with a gap of $C - R = 0.0766$, and the gap is explained quantitatively by Eq. (9): the automatable controls still carry a one-day collection cadence plus a seven-day remediation, an eight-day residual against a periodic latency that on automatable controls averages 272 days before remediation, so the residual fraction $\rho = (C - R)/C = 0.119$ is small and positive, exactly as the derivation requires. The ceiling binds and is not exceeded; the realized reduction sits below it by the residual the automatable controls cannot shed. The diminishing-returns inequality of Eq. (11) is confirmed by the frozen control-days saved, 33,266 against annual and 15,584 against quarterly, with $\Delta(\text{quarterly}) < \Delta(\text{annual})$ by about 17,700 control-days. The concentration relation of Eq. (12) is confirmed by the frozen top-quartile capture of 0.786 ([0.777, 0.796]), well above a uniform-share expectation of one quarter. Each closed-form statement matches the pre-registered numbers it predicts; where the match is approximate rather than exact, as in Eq. (12), it is stated as an approximation and the residual is attributed to the mild variation of the latency factor across controls rather than asserted away.

VII. ROBUSTNESS AND SENSITIVITY

The headline numbers are means over 25 seeds; this section examines how they move across the manual cadence and across the drift distribution, using only the frozen sweeps, and then asks what an adversary can do with the residual the ceiling leaves in place.

A. Sensitivity to the manual cadence

Table II reports the benefit of continuous monitoring against the two manual cadences agencies actually run. The fractional reduction is nearly flat, from 0.567 against annual to 0.552 against quarterly, because it tracks the automatable-exposure share C and that share is a property of the catalog rather than of the cadence. The absolute control-days saved, by contrast, fall by more than half, from 33,266 to 15,584, exactly the diminishing-returns relation of Eq. (11): a quarterly baseline has already removed much of the latency that continuous monitoring would otherwise eliminate, so the remaining headroom is smaller. The practical reading is that the fractional case for automation is cadence-independent but the absolute case is strongest where the manual baseline is weakest.

B. Concentration across the drift distribution

Table III reports how much of the achievable reduction the highest-drift quartile of controls captures, against the

TABLE III

DRIFT CONCENTRATION: REDUCTION CAPTURED BY THE HIGHEST-DRIFT QUARTILE OF CONTROLS (MEAN OVER 25 SEEDS, 95% BCA INTERVAL), AGAINST THE UNIFORM-SHARE BENCHMARK.

Automated subset	Fractional reduction R	Share of full reduction
Top-drift quartile	0.446 [0.435, 0.457]	0.786 [0.777, 0.796]
All automatable controls	0.567 [0.553, 0.581]	1.000
Uniform-share benchmark	n/a	0.250

uniform-share benchmark of one quarter. Automating the top-drift quartile alone realizes a fractional reduction of 0.446 ([0.435, 0.457]), which is 0.786 ([0.777, 0.796]) of the 0.567 reduction available from automating every automatable control. A quarter of the controls thus delivers nearly four-fifths of the benefit, a more than three-fold concentration over the uniform expectation, confirming Eq. (12) on the frozen sweep. The deployment ordering is immediate: instrument the high-drift quartile first.

C. Adversarial robustness of the residual

The ceiling has an adversarial reading that the means conceal. The automatability ceiling guarantees a residual exposure of $1 - C \approx 0.357$ of baseline that no detection speed can lower, and that residual lives entirely in the un-automatable controls, the policy, documentation, and physical controls that a machine cannot check. An adversary who understands the program will not attack the instrumented controls, where a drift is caught within a day; the adversary will concentrate risk in the un-automatable floor, where detection still runs on the periodic cadence T . By moving its activity into the $1 - C$ share, an adversary defeats the automation program without ever tripping a continuous check, because the program has no reach there by construction. The same orthogonal-signal logic that makes a layered defense robust works in reverse here: an automation program that is strong on exactly one axis, machine-checkability, is blind on the complementary axis, and a strategic adversary will choose the blind axis. The residual must therefore be managed by other means, a tighter manual assessment cadence on the specific high-impact policy and physical controls, independent attestation, or human review keyed to change windows, rather than treated as acceptable because the automatable controls are well covered. The ceiling is not only a limit on the benefit; it is a map of where an adversary will go.

VIII. DISCUSSION

The two benefits of compliance as code point in different directions, and conflating them produces the optimism that this study is designed to puncture. Detection latency on automatable controls collapses by a factor of 272; an agency that measures success by how fast it catches a disabled audit log will see a transformative improvement. But the quantity that an authorization decision turns on is exposure, the total control-days a system spends out of compliance, and exposure is bounded by automatability. A reduction of 0.567 against an annual baseline is large and worth pursuing, but it is

closer to half than to the near-elimination that latency figures suggest, and it cannot be pushed past the 0.643 ceiling by any improvement in collection speed. The right program metric is the automatable-exposure share, not the detection-latency reduction.

The concentration result turns this into an actionable ordering. Because the top-drift quartile captures nearly four-fifths of the achievable reduction, an agency should instrument its high-drift, high-automatability technical controls first, the logging, certificate, configuration, and access controls, and expect rapidly diminishing marginal returns thereafter. The diminishing-returns result against cadence sharpens the business case: the gain from automation is largest for an organization still assessing annually and roughly halves for one already on a quarterly cycle, so the investment is most justified where the manual baseline is weakest.

Three deployment rules follow directly. First, size expectations to the automatable-exposure share, not to detection latency, when projecting the compliance-exposure benefit of an automation program. Second, sequence the rollout by drift frequency, because the high-drift quartile dominates the return. Third, treat the un-automatable remainder as an irreducible residual to be managed by other means, such as a tighter manual cadence on the specific high-impact policy and physical controls, rather than as a target for automation.

IX. THREATS TO VALIDITY

Construct validity. The model abstracts compliance drift as a Poisson process with a fixed remediation time. Real drift may be bursty, correlated across controls (a single misconfigured deployment can trip several controls at once), or seasonal around change windows, and remediation time varies with the control and the team. These would shift absolute exposure values but not the structural relationships: an automatability ceiling exists whenever any control is un-automatable, and concentration follows from any heavy-tailed drift-rate distribution.

Internal validity. Because the study is pre-registered and the evaluation seeds were inspected only once, the reported numbers are not the product of analytic search. The development of the model used separate seeds, and no distribution was tuned to clear a threshold. The mechanical evaluation script computes each verdict from the frozen data with no hand-set value.

External validity. The catalog size, the drift-rate mixture, and the automatability probabilities are documented priors, not measurements from any agency, so the absolute magnitudes (272 days, 33,266 control-days) should be read as illustrative of the model rather than as field estimates. The comparative findings, the ceiling, the diminishing returns, and the concentration, depend on the qualitative structure of the model rather than on the specific constants, and we expect them to transfer. Calibrating the drift and automatability distributions against a real control catalog and assessment record is the natural next step and would convert the illustrative magnitudes into estimates.

Statistical validity. Intervals are BCa bootstrap intervals over 25 seeds, appropriate for the small-sample, possibly skewed

statistics reported here; we use interval non-overlap rather than significance testing throughout, avoiding the multiple-comparison pitfalls of repeated p -values.

X. CONCLUSION

Continuous automated evidence collection is transformative for detection latency on the controls it can reach, turning months into a day, but the program-level benefit is governed by two facts that the latency figure hides. A substantial share of controls cannot be machine-checked and sets a hard floor on exposure, and exposure is dominated by the high-drift technical controls, which are also the most automatable. In a pre-registered simulation, continuous monitoring cut detection latency 272-fold on automatable controls, removed 0.567 of annual exposure against a binding ceiling of 0.643, showed diminishing returns as the manual cadence tightened, and concentrated nearly four-fifths of its value in the highest-drift quartile of controls. The deployment guidance is correspondingly concrete: instrument the high-drift automatable controls first, size expectations to the automatable-exposure share, and manage the un-automatable remainder as a residual rather than a target. Calibrating the model against a real control catalog is the next step toward turning these structural findings into field estimates.

DATA AVAILABILITY STATEMENT

The data and code that support the findings of this study are openly available in the research-papers repository at <https://github.com/Harshavardhanmalla-Labs/research-papers/tree/main/paper12>. The manuscript sources, figures, analysis code, and frozen result tables are included and regenerate every reported number deterministically from the fixed seeds.

DISCLOSURE STATEMENT

No potential conflict of interest was reported by the author.

FUNDING

No funding was received.

REFERENCES

- [1] Joint Task Force, "Security and privacy controls for information systems and organizations," National Institute of Standards and Technology, Tech. Rep. NIST Special Publication 800-53 Rev. 5, 2020.
- [2] K. Dempsey, N. S. Chawla, A. Johnson, R. Johnston, A. C. Jones, A. Orebaugh, M. Scholl, and K. Stine, "Information security continuous monitoring (iscm) for federal information systems and organizations," National Institute of Standards and Technology, Tech. Rep. NIST Special Publication 800-137, 2011.
- [3] K. Dempsey, V. Pillitteri, C. Baer, R. Niemeyer, R. Rudman, and S. Urban, "Assessing information security continuous monitoring (iscm) programs: Developing an iscm program assessment," National Institute of Standards and Technology, Tech. Rep. NIST Special Publication 800-137A, 2020.
- [4] National Institute of Standards and Technology, "Open security controls assessment language (OSCAL)," <https://pages.nist.gov/OSCAL/>, 2026, NIST CSRC project; accessed 2026-06-20.
- [5] Open Policy Agent Project (Cloud Native Computing Foundation), "Open policy agent: Policy-based control for cloud native environments," <https://www.openpolicyagent.org/docs>, 2026, graduated CNCF project; accessed 2026-06-20.

- [6] Microsoft, “Powershell desired state configuration (DSC),” <https://learn.microsoft.com/powershell/dsc/overview>, 2026, accessed 2026-06-20.
- [7] B. Efron, “Better bootstrap confidence intervals,” *Journal of the American Statistical Association*, vol. 82, no. 397, pp. 171-185, 1987.
- [8] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. Chapman & Hall/CRC, 1994.
- [9] C. Wohlin, P. Runeson, M. Höst, M. C. Ohlsson, B. Regnell, and A. Wesslén, *Experimentation in Software Engineering*. Springer, 2012.
- [10] National Institute of Standards and Technology, “Standards for security categorization of federal information and information systems,” U.S. Department of Commerce, Tech. Rep. FIPS PUB 199, 2004.
- [11] K. Stine, R. Kissel, W. C. Barker, J. Fahlsing, and J. Gulick, “Guide for mapping types of information and information systems to security categories,” National Institute of Standards and Technology, Tech. Rep. NIST Special Publication 800-60 Rev. 1, 2008.
- [12] A. Rahman, C. Parnin, and L. Williams, “The seven sins: Security smells in infrastructure as code scripts,” in *Proceedings of the 41st International Conference on Software Engineering (ICSE)*, 2019, pp. 164-175.
- [13] A. Rahman and L. Williams, “Different kind of smells: Security smells in infrastructure as code scripts,” *IEEE Security & Privacy*, vol. 19, no. 3, pp. 33-41, 2021.
- [14] J. O. Kephart and D. M. Chess, “The vision of autonomic computing,” *IEEE Computer*, vol. 36, no. 1, pp. 41-50, 2003.
- [15] National Institute of Standards and Technology, “Guide to enterprise patch management planning: Preventive maintenance for technology,” National Institute of Standards and Technology, Tech. Rep. NIST Special Publication 800-40 Rev. 4, 2022.
- [16] Cybersecurity and Infrastructure Security Agency (CISA), “Binding operational directive 22-01: Reducing the significant risk of known exploited vulnerabilities,” U.S. Department of Homeland Security, Tech. Rep., Nov. 2021, <https://www.cisa.gov/news-events/directives/bod-22-01-reducing-significant-risk-known-exploited-vulnerabilities>.
- [17] Cybersecurity and Infrastructure Security Agency (CISA), “Binding operational directive 23-01: Improving asset visibility and vulnerability detection on federal networks,” U.S. Department of Homeland Security, Tech. Rep., Oct. 2022, <https://www.cisa.gov/news-events/directives/bod-23-01-improving-asset-visibility-and-vulnerability-detection-federal-networks>.
- [18] Verizon, “2026 data breach investigations report (DBIR),” Verizon Business, Tech. Rep., 2026, <https://www.verizon.com/business/resources/reports/dbir/>.
- [19] H. Malla, “HygienePrio: Cyber-hygiene signal augmentation for EPSS-weighted vulnerability prioritization,” <https://github.com/Harshavardhanmalla-Labs/research-papers/tree/main/paper4>, 2026, companion paper; source, code, and frozen results in the same repository.
- [20] H. Malla, “Context-aware vulnerability prioritization for government endpoint fleets,” <https://github.com/Harshavardhanmalla-Labs/research-papers/tree/main/paper1>, 2026, companion paper; source, code, and frozen results in the same repository.
- [21] H. Malla, “Policy-as-code for preventive compliance enforcement: prevention versus detection for recurring control violations, with a CJIS case study,” <https://github.com/Harshavardhanmalla-Labs/research-papers/tree/main/paper13>, 2026, companion paper.
- [22] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed. Wiley, 2006.